VOICE AI RECEPTIONIST

Build a $5K/Month Solution in 3 Days (Even If You Can't Code)

By Josh | Amalfi AI

I BUILT THIS IN 3 DAYS WITH NO CODING BACKGROUND

Let me start with the uncomfortable truth: I'm not a developer. Never took a CS course. Can't write a for-loop from memory.

Yet I built a voice AI receptionist that:

• Answers calls 24/7
• Qualifies leads with custom questions
• Books appointments directly into calendars
• Sends follow-up messages automatically
• Handles multiple calls simultaneously

Build time: 3 days.
What I charge: $5K/month per client.
Current clients: Running this for 2 businesses, $10K MRR from this product alone.

How? **I let Claude build it for me.**

This guide is going to show you exactly how to build this system. Not theory. Not concepts. The actual step-by-step technical process, with every tool, every prompt, every integration.

By the end, you'll have a working voice AI receptionist you can demo to prospects. Then it's just sales.

WHY THIS PRODUCT SELLS ITSELF

Before we build, understand why this is money:

THE MATH THAT CLOSES DEALS

Scenario 1: Medical Practice

• Receptionist salary: $3,000/month
• Works: 8am-5pm, Monday-Friday
• Misses: ~25% of calls (bathroom, lunch, busy)
• Each missed call: Potential $200-500 appointment

Monthly cost: $3K salary + ~$2-3K in lost bookings = $5-6K total pain

Your solution:
• Never misses calls (24/7 availability)
• Books appointments instantly
• Costs: $5K/month
• Saves them: $3K salary + stops revenue loss

ROI: They break even immediately, gain after-hours coverage for free.

Scenario 2: Real Estate Agency

• 50+ inbound calls weekly
• Receptionist handles during office hours only
• After-hours/weekend calls go to voicemail
• Response time: 24-48 hours
• Conversion rate drops 60% with delayed response

Your solution:
• Answers instantly, 24/7
• Qualifies buyer intent
• Books viewings automatically
• Costs: $4K/month

One extra deal per month = $5-10K commission. Your fee pays for itself 2-3X.

WHO NEEDS THIS (TARGET INDUSTRIES)

This product crushes in:

• **Medical practices:** Appointments, patient screening
• **Law firms:** Case intake, consultation booking

- **Real estate:** Lead qualification, viewing bookings
- **Home services:** Quote requests, job scheduling
- **Fitness/wellness:** Class bookings, membership inquiries
- **Salons/spas:** Appointment scheduling, service questions

Pattern: **High call volume + appointment-based business = perfect fit.**

THE COMPLETE TECH STACK (WHAT YOU ACTUALLY NEED)

Here's every tool you need and what it costs:

CORE COMPONENTS:

1. Vapi.ai - Voice AI Platform

What it does: Handles phone calls, voice recognition, speech synthesis

Cost: $0.05-0.12 per minute of call time

Why this one: Best voice quality, easiest setup, most reliable

2. ElevenLabs - Voice Cloning (Optional but Recommended)

What it does: Natural-sounding voice synthesis

Cost: $22/month for professional tier

Why: Makes the AI sound human, not robotic

3. Claude API (Anthropic)

What it does: Powers the conversational intelligence

Cost: ~$0.50-2 per call (depends on length)

Why: Best at understanding context, following instructions

4. Make.com (formerly Integromat)

What it does: Connects everything together (calendar, CRM, notifications)

Cost: $9-29/month depending on usage

Why: Visual workflow builder, no coding needed

5. Google Calendar API (or client's calendar)

What it does: Books appointments directly

Cost: Free

Why: Everyone uses it, easy integration

TOTAL MONTHLY COST: $50-100 depending on call volume
What you charge: $4-5K/month
Your margin: 95%+

ALTERNATIVE STACKS (If Budget Constrained):

• Replace Vapi with Twilio + custom setup (harder, cheaper)
• Replace ElevenLabs with Google TTS (robotic but functional)
• Replace Make.com with Zapier (similar pricing, less flexible)
• Replace Claude with GPT-4 (slightly cheaper, less reliable context)

My recommendation: **Use the premium stack.** You're charging $5K/month. Spend the $100 to make it flawless.

THE 3-DAY BUILD PROCESS

I'm giving you the exact sequence. Follow it precisely.

DAY 1: FOUNDATION & BASIC CALL HANDLING

Hour 1-2: Vapi Setup

1. Sign up at vapi.ai
2. Create new assistant
3. Configure base settings:

• Voice: Choose from Vapi's library (I use 'Sarah' - professional female voice)
• Model: Claude Sonnet (select from dropdown)
• First message: "Thanks for calling [Business Name]. How can I help you today?"

Hour 3-4: Writing the System Prompt

This is critical. The system prompt controls everything the AI says and does.

Here's the framework I use (customize for client):

"You are the receptionist for [Business Name], a [industry] business. Your job is to:

1. Greet callers professionally and warmly
2. Determine their needs by asking: [key qualifying question]
3. If they want to book: Collect name, phone, preferred date/time
4. If they have questions: Answer using this info: [provide business details]
5. End calls by confirming next steps

Rules:
- Keep responses under 30 seconds
- Never say you're AI unless asked directly
- If you don't know something, say you'll have someone call back
- Transfer to human for: [emergency situations, complaints, complex issues]

Tone: Professional, friendly, helpful. Like a receptionist who genuinely cares."

Hour 5-6: Test Basic Calling

Vapi gives you a test number. Call it 20+ times:

• Test different scenarios (booking, questions, complaints)
• Note where it fails or sounds awkward
• Refine the system prompt
• Adjust voice speed/tone settings

This is tedious. Do it anyway. Every awkward pause you catch now = one less prospect thinking "this sounds robotic."

Hour 7-8: Add Conversation Tools

In Vapi, add 'Tools' (functions the AI can call):

Tool 1: Check Calendar Availability
• Connects to Google Calendar
• Returns available time slots
• AI uses this to offer options

Tool 2: Book Appointment
• Takes: name, phone, date, time, notes
• Creates calendar event
• Sends confirmation SMS

Vapi has templates for both. Use them. Customize the calendar connection.

END OF DAY 1: You have a working call handler that can answer and book appointments.

DAY 2: INTELLIGENCE & AUTOMATION

Hour 1-3: Lead Qualification Logic

Now make it smart. Add qualification questions to the system prompt:

"When someone calls about [service], ask these questions in order:
1. [Qualifying question 1 - budget/timeline]
2. [Qualifying question 2 - specific need]
3. [Qualifying question 3 - decision maker]

Based on answers, categorize as:
- HOT: Ready to book, good fit
- WARM: Interested but needs follow-up
- COLD: Not a fit or tire-kicker

Tag each lead in the system with this category."

Test this extensively. The qualification saves your client hours of wasted time.

Hour 4-6: Make.com Integration Setup

Open Make.com. Create new scenario. Add these modules:

Trigger: Webhook (Vapi sends call data here)
Action 1: Create Google Calendar event
Action 2: Send SMS confirmation (Twilio)
Action 3: Log to Google Sheets (call tracking)
Action 4: Send Slack notification to client (optional)

Connect Vapi to this webhook. Every call now triggers this automation.

Hour 7-8: Build the Follow-Up System

For WARM leads (interested but didn't book), set up:

• 2-hour follow-up SMS: "Hi [Name], this is [Business]. We spoke earlier about [topic]. Still interested? Reply YES to book or call us back."
• 24-hour follow-up email (if provided)
• 3-day final touchpoint

This is where most solutions fail. Yours won't. You're capturing leads others lose.

END OF DAY 2: You have an intelligent system that qualifies, books, and follows up automatically.

DAY 3: POLISH & CLIENT-READY DEPLOYMENT

Hour 1-2: Voice Quality Optimization

If using ElevenLabs (recommended):

• Clone a professional voice (use Fiverr to get voice samples if needed)
• Or choose from ElevenLabs library
• Integrate into Vapi (they have native connection)
• Test calls - should sound 90% human

This $22/month subscription is the difference between "cool demo" and "I can't tell this isn't human."

Hour 3-4: Handle Edge Cases

Test failure scenarios:

• What if caller speaks another language?
• What if they're angry/upset?
• What if they need emergency help?
• What if calendar is fully booked?
• What if they request something outside scope?

Add handling to system prompt:

"If caller is emotional/upset: 'I understand this is frustrating. Let me have [Owner] call you directly within 2 hours. Can I have your best number?'

If calendar full: 'We're fully booked this week. I can get you on the waitlist or offer [alternative date]. What works better?'"

Hour 5-6: Build the Client Dashboard

Your client needs visibility. Create:

Google Sheet Dashboard with:
• Call log (date, time, caller, outcome)
• Lead quality breakdown (HOT/WARM/COLD count)
• Booked appointments (auto-populated from calendar)
• Response time metrics
• Monthly summary stats

This updates in real-time via Make.com. Client sees value constantly.

Hour 7-8: Final Testing & Documentation

• Call the system 50+ times (yes, really)

• Have friends/family call and try to break it

• Document any issues and fix them

• Write a 1-page guide for client on how to update hours/services/pricing

END OF DAY 3: You have a production-ready voice AI receptionist worth $5K/month.

THE EXACT CLAUDE PROMPTS I USED

Here's how I built this without coding. These are the actual prompts:

Prompt 1: System Design

"I need to build a voice AI receptionist for [industry] businesses. It needs to:
1. Answer calls 24/7
2. Ask qualifying questions
3. Book appointments into Google Calendar
4. Send SMS confirmations
5. Log everything to a dashboard

What's the best tech stack? Give me specific tools and explain why each one."

Prompt 2: System Prompt Writing

"I'm using Vapi.ai with Claude as the model. I need a system prompt for a receptionist at [Business Name]. They need to:
- Greet callers professionally
- Ask [these qualifying questions]
- Book appointments if qualified
- Handle objections about price/availability
- Sound like a real person, not AI

Write me the complete system prompt. Include tone guidelines and edge case handling."

Prompt 3: Make.com Workflow

"I need a Make.com scenario that:
1. Receives webhook data from Vapi (call transcript, caller info, booking details)
2. Creates Google Calendar event
3. Sends SMS via Twilio
4. Logs to Google Sheets
5. Sends Slack notification

Walk me through the exact modules I need and how to connect them. I'm not technical so explain each step."

Prompt 4: Troubleshooting

"The AI keeps [specific problem]. Here's my system prompt: [paste it]

What's wrong and how do I fix it?"

I used this prompt at least 20 times during the build. Claude debugs faster than any human.

HOW TO DEMO THIS TO CLOSE DEALS

You built it. Now sell it. Here's the demo process:

PRE-DEMO SETUP (Critical)

Before the discovery call, customize the demo for THEIR business:

• Change business name in greeting
• Update services/pricing in system prompt
• Adjust qualifying questions to their industry
• Pre-populate their calendar availability

This takes 15 minutes. The impact is MASSIVE. They hear THEIR business, not a generic demo.

THE DEMO SCRIPT

"Before I show you, quick context: Your receptionist costs $3K monthly and works 40 hours/week. This system never sleeps, never takes breaks, costs less, and captures every lead. Watch."

Then CALL IT LIVE during the meeting. Speaker mode on.

Call 1 - Simple Booking:
"Hi, I need to schedule an appointment for [service]."
AI handles it, books it, confirms details.

Call 2 - Qualifying Lead:
"I'm interested but need to know prices first."
AI answers pricing questions, qualifies budget, books or follows up.

Call 3 - After Hours:
"It's 10 PM. Let me show you what happens..."
AI answers perfectly. Point out: "Your human receptionist is home. This one just caught a lead."

Then pull up the dashboard:

"Here's where all the calls log in real-time. See the lead quality tags? You'll know immediately which calls to prioritize. Everything auto-books to your calendar."

HANDLING OBJECTIONS

"Sounds robotic"

Response: Call it again. Have THEM test it. Let them hear the natural voice. Then: "What did it miss?"

"What if it makes mistakes?"

Response: "It will. Every receptionist does. Difference? I fix it in 10 minutes by updating the prompt. Training a human takes days."

"We need a human touch"

Response: "Agreed. This handles the routine stuff - bookings, basic questions. Your team focuses on the calls that need human touch - complaints, complex issues, relationship-building."

"Too expensive"

Response: "Your receptionist is $3K and works 160 hours monthly. This is $5K and works 720 hours monthly. Per-hour cost is 70% cheaper AND it never misses after-hours calls. What's expensive is losing leads."

DEPLOYING FOR CLIENTS (THE 21-DAY DELIVERY)

You closed the deal. Now deliver. Here's the timeline:

WEEK 1: Setup & Customization

• Day 1-2: Discovery call (record with Read.ai), feed to Claude for analysis
• Day 3-4: Build customized version (business details, services, pricing, tone)
• Day 5: Internal testing (call it 100+ times)
• Day 6-7: Client review call, make adjustments

WEEK 2: Testing & Training

• Day 8-10: Client tests extensively, provides feedback
• Day 11-12: Refinements based on feedback
• Day 13-14: Train client's team on dashboard, how to handle escalations

WEEK 3: Launch & Monitoring

• Day 15: Go live with real phone number
• Day 16-18: Monitor every call, fix issues immediately
• Day 19-20: Optimization based on real caller behavior
• Day 21: Handoff meeting, confirm everything works

POST-LAUNCH SUPPORT

Include 30 days of adjustment period:

• Update system prompt for edge cases they discover
• Adjust qualification questions based on actual calls
• Optimize booking flow if drop-offs happen
• Add new services/pricing as they evolve

After 30 days, charge $100-200 per significant change. Or include unlimited adjustments in monthly fee.

SCALING BEYOND CLIENT #1

Once you have one working version, scaling is copy-paste:

THE TEMPLATE APPROACH

After client #1, you have:

• Base system prompt (85% reusable)
• Make.com workflow (100% reusable)
• Dashboard template (100% reusable)
• Demo process (100% reusable)
• Delivery timeline (100% reusable)

Client #2 build time: 1 day (not 3).
Client #3 build time: 4 hours.
Client #4+: 2-3 hours.

You're copy-pasting and customizing, not building from scratch.

INDUSTRY-SPECIFIC VERSIONS

Build variations for different verticals:

Medical Practice Version:
• Handles insurance verification
• Collects patient symptoms (pre-appointment screening)
• HIPAA-compliant logging
• Emergency escalation triggers

Real Estate Version:
• Qualifies buyer budget
• Schedules property viewings
• Sends listing details via SMS
• Tracks lead source for attribution

Home Services Version:
• Collects job details (photos, dimensions, scope)
• Provides quote ranges
• Books site visits
• Sends confirmation with tech profile

Each version = new market you can sell into.

PRICING AS YOU SCALE

Client #1-2: $4-5K/month (you're learning, charge less)
Client #3-5: $5-6K/month (you're fast, charge more)
Client #6+: $6-8K/month (you're an expert, premium pricing)

Or introduce tiers:

Basic: $3K/month - Call answering + booking only
Pro: $5K/month - + Lead qualification + follow-up
Enterprise: $8K/month - + CRM integration + custom reporting

THE PART NOBODY TALKS ABOUT

I've given you everything. The tech. The prompts. The demo script. The delivery timeline.

But here's what I can't give you:

The willingness to iterate.

This won't work perfectly on day 1. The AI will sound awkward. Calls will drop. Bookings will fail.

That's normal. Expected. Part of the process.

I called my first system 200+ times before it sounded natural. I rewrote the system prompt 30+ times. I rebuilt the Make.com workflow 4 times.

Most people quit after 10 test calls: "This doesn't work."

Winners iterate: "What's not working? How do I fix it?"

The willingness to learn while building.

You will get stuck. Vapi will throw errors you don't understand. Make.com will break mysteriously. Claude will misinterpret your prompts.

When that happens, you have two options:

1. Give up: "This is too technical for me"
2. Google it, ask Claude for help, watch a YouTube video, figure it out

I chose #2 at least 50 times during my first build. Every time I got unstuck, I got smarter.

By client #3, nothing breaks me. Because I've already solved every problem once.

The courage to sell before it's perfect.

Your first version will have flaws. You'll notice them. You'll want to fix everything before demoing.

Don't. Demo at 80%. Prospects don't care about perfection. They care about "is this better than what I have now?"

My first demo had:

• Slightly robotic voice (before ElevenLabs)
• Occasional awkward pauses
• Basic dashboard (just a Google Sheet)

Still closed the deal. Because it was better than their $3K receptionist who missed calls.

WANT THE COMPLETE SYSTEM PROMPTS AND BUILD TEMPLATES?

This guide gave you the roadmap. But when you're staring at a blank Vapi screen at midnight... when your Make.com workflow breaks and you don't know why... when a prospect asks a question you didn't prepare for...

That's when you need the templates.

Join the Amalfi AI Mentorship: $97 one-time

• Complete system prompts for 5 industries (copy-paste ready)

• Make.com workflow templates (import and customize)

• Dashboard templates (Google Sheets + visualization)

• Demo script with objection handling

• Client onboarding checklist and delivery timeline

• Troubleshooting guide for common issues

• Pricing calculator (what to charge based on client size)

• Private community with other builders solving the same problems

Optional: $27/month for weekly build sessions and live troubleshooting.

Stop theorizing. Start building. Stop watching demos. Start creating them.

www.amalfiai.com/mentorship